

1 grepq: A Rust application that quickly filters
2 FASTQ files by matching sequences to a set
3 of regular expressions

4 Nicholas D. Crosbie

5 25 January 2025

6 **grepq: A Rust application that quickly filters**
7 **FASTQ files by matching sequences to a set of**
8 **regular expressions**

9 *Nicholas D. Crosbie, Melbourne Veterinary School, University of Melbourne,*
10 *Parkville, Victoria, Australia*

11 ORCID: 0000-0002-0319-4248

12 January 15, 2025

13 **Keywords:** bioinformatics, FASTQ records, regular expressions,
14 Rust, variants

15 **Abstract**

16 Regular expressions (regex) (Kleene 1951) have been an im-
17 portant tool for finding patterns in biological codes for decades
18 (Hodgman 2000 and citations therein), and unlike fuzzy-finding
19 approaches, do not result in approximate matches. The perfor-
20 mance of regular expressions can be slow, however, especially
21 when searching for matching patterns in large files. *grepq* is
22 a Rust application that quickly filters FASTQ files by matching
23 sequences to a set of regular expressions. *grepq* is designed
24 with a focus on performance and scalability, is easy to install and
25 easy to use, enabling users to quickly filter large FASTQ files, to
26 enumerate variants and update the order in which patterns are
27 matched against sequences through an in-built *tune* command.
28 *grepq* is open-source and available on *GitHub* and *Crates.io*.

29 **Statement of need**

30 The ability to quickly filter FASTQ files by matching sequences to
31 a set of regular expressions is an important task in bioinformatics,
32 especially when working with large datasets. The importance and
33 challenge of this task will only grow as sequencing technologies
34 continue to advance and produce ever larger datasets (Katz et
35 al. 2022). The uses cases of *grepq* are diverse, and include pre-

36 processing of FASTQ files before downstream analysis, quality
37 control of sequencing data, and filtering out unwanted sequences.
38 Where decisions need be made quickly, such as in a clinical set-
39 tings (Bachurin et al. 2024), biosecurity (Valdivia-Granda 2012),
40 and wastewater-based epidemiology in support of public health
41 measures (Choi et al. 2018; Sims and Kasprzyk-Hordern 2020;
42 Xylogiannopoulos 2021; Merrett et al. 2024), the ability to quickly
43 filter FASTQ files and enumerate variants by matching sequences
44 to a set of regular expressions is attractive as it circumvents the
45 need for more time-consuming bioinformatic workflows.

46 Regular expressions are a powerful tool for matching sequences,
47 but they can be slow and inefficient when working with large
48 datasets. Furthermore, general purpose tools like *grep* (Free
49 Software Foundation 2023) and *ripgrep* (A. Gallant 2025) are
50 not optimized for the specific task of filtering FASTQ files, and
51 occasionally yield false positives as they scan the entire FASTQ
52 record, including the sequence quality field. Tools such *awk* (Aho,
53 Kernighan, and Weinberger 1988) and *gawk* (Free Software
54 Foundation 2024) can be used to filter FASTQ files without yield-
55 ing false positives, but they are significantly slower than *grepq*
56 and can require the development of more complex scripts to
57 achieve the same result.

58 Implementation

59 *grepq* is implemented in Rust, a systems programming language
60 known for its safety features, which help prevent common pro-
61 gramming errors such as null pointer dereferences and buffer over-
62 flows. These features make Rust an ideal choice for implementing
63 a tool like *grepq*, which needs to be fast, efficient, and reliable.

64 Furthermore, *grepq* obtains its performance and reliability, in part,
65 by using the *seq_io* (Schlegel and Seyboldt 2025) and *regex* (Gal-
66 lant et al. 2025b) libraries. The *seq_io* library is a well-tested
67 library for parsing FASTQ files, designed to be fast and efficient,
68 and which includes a module for parallel processing of FASTQ
69 records through multi-threading. The *regex* library is designed
70 to work with regular expressions and sets of regular expressions,
71 and is known to be one of the fastest regular expression libraries
72 currently available (Gallant et al. 2025a). The *regex* library sup-
73 ports Perl-like regular expressions without look-around or backref-
74 erences (documented at https://docs.rs/regex/1.*/regex/#syntax).

75 Further performance gains were obtained by:

- 76 • use of the *RegexSet* struct from the *regex* library to match
77 multiple regular expressions against a sequence in a single
78 pass, rather than matching each regular expression individu-
79 ally (the *RegexSet* is created and compiled once before en-

80 tering any loop that processes the FASTQ records, avoiding
81 the overhead of recompiling the regular expressions for each
82 record)

- 83 • multi-threading to process the records within an input FASTQ
84 file in parallel through use of multiple CPU cores
- 85 • use of the *zlib-ng* backend to the *flate2* library to read and
86 write gzip-compressed FASTQ files, which is faster than the
87 default *miniz_oxide* backend
- 88 • use of an optimised global memory allocator (the *mimalloc*
89 library (Mutiple, n.d.)) to reduce memory fragmentation and
90 improve memory allocation and deallocation performance
- 91 • buffer reuse to reduce the number of memory allocations and
92 deallocations
- 93 • use of byte slices to avoid the overhead of converting to and
94 from string types
- 95 • in-lining of performance-critical functions
- 96 • use of the *write_all* I/O operation that ensures the data is writ-
97 ten in one go, rather than writing data in smaller chunks

98 **Feature set**

99 *grepq* has the following features:

- 100 • support for presence and absence (inverted) matching of a
101 set of regular expressions

- 102 • IUPAC ambiguity code support (N, R, Y, etc.)
- 103 • support for gzip and zstd compression (reading and writing)
- 104 • JSON support for pattern file input and *tune* command output,
- 105 allowing named regular expression sets and named regular
- 106 expressions (pattern files can also be in plain text)
- 107 • the ability to set predicates to filter FASTQ records on the
- 108 header field (= record ID line) using a regular expression, min-
- 109 imum sequence length, and minimum average quality score
- 110 (supports Phred+33 and Phred+64)
- 111 • the ability to output matched sequences to one of four formats
- 112 (including FASTQ and FASTA)
- 113 • the ability to tune the pattern file and enumerate variants with
- 114 the *tune* command: this command will output a plain text or
- 115 JSON file with the patterns sorted by their frequency of oc-
- 116 currence in the input FASTQ file or gzip-compressed FASTQ
- 117 file (or a user-specified number of FASTQ records). This can
- 118 be useful for optimizing the pattern file for performance, for
- 119 example by removing patterns that are rarely matched and
- 120 reordering nucleotides within the variable regions of the pat-
- 121 terns to improve matching efficiency
- 122 • the ability to count and summarise the total number of records
- 123 and the number of matching records (or records that don't
- 124 match in the case of inverted matching) in the input FASTQ
- 125 file

126 Other than when the *tune* command is run, a FASTQ record is
127 deemed to match (and hence provided in the output) when any
128 of the regular expressions in the pattern file match the sequence
129 field of the FASTQ record. Example output of the *tune* command
130 (when given with the **-json-matches** flag) is shown below:

```
# For each matched pattern in a search of the first  
# 20000 records of a gzip-compressed FASTQ file, print  
# the pattern and the number of matches to a JSON file  
# called matches.json, and include the top three most  
# frequent variants of each pattern, and their respective  
# counts
```

```
grepq --read-gzip 16S-no-iupac.json SRX26365298.fastq.gz \  
tune -n 20000 -c --names --json-matches --variants 3
```

131 Output (abridged) written to matches.json:

```
{  
  "regexSet": {  
    "regex": [  
      {  
        "mostFrequentVariants": [  
          {  
            "count": 219,  
            "variant": "GAATTGACGGGG"          }  
        ]  
      }  
    ]  
  }  
}
```

```

    },
    {
        "count": 43,
        "variant": "AAATTGACGGGG"
    },
    {
        "count": 21,
        "variant": "GAATTGGCGGGG"
    }
],
"regexCount": 287,
"regexName": "Primer contig 06a",
"regexString": "[AG]AAT[AT]G[AG]CGGGG"
},
{
    "mostFrequentVariants": [
        {
            "count": 221,
            "variant": "CCCCGTCAATTC"
        },
        {
            "count": 43,
            "variant": "CCCCGTCAATTT"
        }
    ],

```



```

        {
            "count": 25,
            "variant": "CCCGCCAATTC"
        }
    ],
    "regexCount": 298,
    "regexName": "Primer contig 06aR",
    "regexString": "CCCG[CT]C[AT]ATT[CT] "
}

],
    "regexSetName": "conserved 16S rRNA regions"
}

```

132 When the count option (-c) is given with the *tune* command, *grepq*
 133 will count the number of FASTQ records containing a sequence
 134 that is matched, for each matching regular expression in the pat-
 135 tern file. If, however, there are multiple occurrences of a given
 136 regular expression within a FASTQ record sequence field, *grepq*
 137 will count this as one match. When the count option (-c) is not
 138 given with the *tune* command, *grepq* provides the total number of
 139 matching FASTQ records for the set of regular expressions in the
 140 pattern file.

141 Colorized output for matching regular expressions is not imple-

mented to maximise speed and minimise code complexity, but can be achieved by piping the output to *grep* or *ripgrep* for testing purposes.

Performance

The performance of *grepq* was compared to that of *fqgrep*, *seqkit grep*, *ripgrep*, *grep*, *awk*, and *gawk* using the benchmarking tool *hyperfine*. The test conditions and results are shown in **Table 1**, **Table 2** and **Table 3**.

Table 1: Wall times and speedup of various tools for filtering FASTQ records against a set of regular expressions. Test FASTQ file: SRX26365298.fastq (uncompressed) was 874MB in size, and contained 869,034 records.

tool	wall time (s)		speedup		
	mean	S.D.	× <i>grep</i>	× <i>ripgrep</i>	× <i>awk</i>
<i>grepq</i>	0.192	0.010	1796.76	18.62	863.52
<i>fqgrep</i>	0.338	0.005	1017.61	10.55	489.07
<i>ripgrep</i>	3.568	0.005	96.49	1.00	46.37
<i>seqkit grep</i>	2.885	0.011	119.33	1.24	57.35
<i>grep</i>	344.259	0.545	1.00	0.01	0.48
<i>awk</i>	165.451	1.590	2.08	0.02	1.00
<i>gawk</i>	287.662	1.682	1.20	0.01	0.58

grepq v1.4.0, *fqgrep* v.1.02, *ripgrep* v14.1.1, *seqkit grep* v.2.9.0, *grep* 2.6.0-FreeBSD, *awk* v. 20200816, and *gawk* v.5.3.1. *fqgrep* and *seqkit grep* were run with default settings, *ripgrep* was run with **-B 1 -A 2 --colors 'match:none' --no-line-number**, and *grep* was run with **-B 1 -A 2 --color=never**. *awk* and *gawk* scripts were also configured to output matching records in FASTQ format. The pattern file contained 30 regular expression representing the 12-mers (and

158 their reverse complement) from Table 3 of Martinez-Porchas et al. (2017). The wall times, given in
 159 seconds, are the mean of 10 runs, and S.D. is the standard deviation of the wall times, also given
 160 in seconds.

161 **Table 2:** Wall times and speedup of various tools for filtering gzip-compressed
 162 FASTQ records against a set of regular expressions. Test FASTQ file:
 163 SRX26365298.fastq.gz was 266MB in size, and contained 869,034 records.

tool	wall time (s)		speedup
	mean	S.D.	× ripgrep
<i>grepq</i>	1.703	0.002	2.10
<i>fqgrep</i>	1.834	0.005	1.95
<i>ripgrep</i>	3.584	0.013	1.00

164 Test conditions and tool versions as above, but *grepq* was run with the **–read-gzip** option, *fqgrep*
 165 with the **-Z** option, and *ripgrep* with the **-z** option. SRX26365298.fastq was gzip-compressed using
 166 the *gzip* v.448.0.3 command (Apple Inc. 2019) using default (level 6) settings. The pattern file
 167 contained 30 regular expression representing the 12-mers (and their reverse complement) from
 168 Table 3 of Martinez-Porchas et al. (2017). The wall times, given in seconds, are the mean of 10
 169 runs, and S.D. is the standard deviation of the wall times, also given in seconds.

Table 3: Wall times and speedup of various tools for filtering FASTQ records against a set of regular expressions. Test FASTQ file: SRX22685872.fastq was 104GB in size, and contained 139,700,067 records.

tool	wall time (s)		speedup
	mean	S.D.	× ripgrep
	Uncompressed		
grepq	26.972	0.244	4.41
fqgrep	50.525	0.501	2.36
ripgrep	119.047	1.227	1.00
	gzip-compressed		
grepq	149.172	1.054	0.98
fqgrep	169.537	0.934	0.86
ripgrep	144.333	0.243	1.00

Test conditions and tool versions as described in the footnote to Table 1. Note that when *grepq* was run on the gzip-compressed file, a memory resident time for the *grepq* process of 116M as reported by the *top* command (Apple Inc. 2023c). *fastq-dump* v3.1.1 (Sherry et al. 2012) was used to download SRX22685872 as a gzip compressed file from the NCBI SRA. The pattern file contained 30 regular expression representing the 12-mers (and their reverse compliment) from Table 3 of Martinez-Porchas et al. (2017). The wall times, given in seconds, are the mean of 10 runs, and S.D. is the standard deviation of the wall times, also given in seconds.

Testing

The output of *grepq* was compared against the output of *fqgrep*, *seqkit grep*, *ripgrep*, *grep*, *awk* and *gawk*, using the *stat* command (Apple Inc. 2023b), and any difference investigated using the *diff* command (Apple Inc. 2023a). Furthermore, a custom utility, *spikeq* (Crosbie 2024b), was developed to generate synthetic FASTQ files with a known number of records and sequences with user-specified lengths that were spiked with a set of regular expressions a

187 known number of times. This utility was used to test the performance of *grepq*
188 and the aforementioned tools under controlled conditions.

189 Finally, a bash test script (see *examples/test.sh*, available at *grepq*'s Github
190 repository) and a simple Rust CLI application, *predate* (Crosbie 2024a), were
191 developed and utilised to automate system testing, and to monitor for perfor-
192 mance regressions.

193 *grepq* has been tested on macOS 15.0.1 (Apple M1 Max) and Linux Ubuntu
194 20.04.6 LTS (AMD EPYC 7763 64-Core Processor). It may work on other plat-
195 forms, but this has not been tested.

196 **Availability and documentation**

197 *grepq* is open-source and available at *GitHub* ([https://github.com/Rbfinch/gre](https://github.com/Rbfinch/grepq)
198 [pq](https://github.com/Rbfinch/grepq)) and *Crates.io* (<https://crates.io/crates/grepq>).

199 Documentation and installation instructions for *grepq* are available at the same
200 Github repository, and through the **-h** and **-help** command-line options, which
201 includes a list of all available commands and options, and examples of how to
202 use them. Example pattern files in plain text and JSON format are also provided,
203 as well as test scripts. *grepq* is distributed under the MIT license.

204 **Conclusion**

205 The performance of *grepq* was compared to that of *fqgrep*, *seqkit grep*, *ripgrep*,
206 *grep*, *awk*, and *gawk* using the benchmarking tool *hyperfine*. For an uncom-
207 pressed FASTQ file 874MB in size, containing 869,034 records, *grepq* was
208 significantly faster than the other tools tested, with a speedup of 1797 times
209 relative to *grep*, 864 times relative to *awk*, and 19 times relative to *ripgrep*. For

210 a larger uncompressed FASTQ file (104GB in size, and containing 139,700,067
211 records), *grepq* was 4.4 times faster than *ripgrep* and marginally slower or of
212 equivalent speed to *ripgrep* where the same large file was gzip-compressed.
213 When coupled with its exceptional runtime performance, *grepq*'s feature set
214 make it a powerful and flexible tool for filtering large FASTQ files.

215 Acknowledgements

216 I'm grateful to my family for their patience and support during the development
217 of *grepq*. I would also like to thank the developers of the *seq_io*, *regex*, *mimalloc*
218 and *flate2* libraries for their excellent work, and the developers of the *hyperfine*
219 benchmarking tool for making it easy to compare the performance of different
220 tools. Finally, I would like to thank the authors of the *ripgrep* and *fqgrep* tools
221 for providing inspiration for *grepq*.

222 Conflicts of interest

223 The author declares no conflicts of interest.

224 References

- 225 Aho, Alfred V., Brian W. Kernighan, and Peter J. Weinberger. 1988. *The AWK*
226 *Programming Language*. <https://www.cs.princeton.edu/~bwk/btl.mirror/>.
- 227 Apple Inc. 2019. *The Gzip Command*.
- 228 ———. 2023a. *The Diff Command*.
- 229 ———. 2023b. *The Stat Command*.
- 230 ———. 2023c. *The Top Command*.
- 231 Bachurin, Stanislav S, Mikhail V Yurushkin, Ilya A Slynko, Mikhail E Kletskii,
232 Oleg N Burov, and Dmitriy P Berezovskiy. 2024. "Structural Peculiarities

233 of Tandem Repeats and Their Clinical Significance.” *Biochemical and Bio-*
 234 *physical Research Communications* 692: 149349.

235 Choi, Phil M, Ben J Tschärke, Erica Donner, Jake W O’Brien, Sharon C Grant,
 236 Sarit L Kaserzon, Rachel Mackie, et al. 2018. “Wastewater-Based Epidemi-
 237 ology Biomarkers: Past, Present and Future.” *TrAC Trends in Analytical*
 238 *Chemistry* 105: 453–69.

239 Crosbie, Nicholas D. 2024a. “predate: Catch bugs and performance regres-
 240 sions through automated system testing.” <https://github.com/Rbfinch/predate>.
 241 te.

242 ———. 2024b. “spikeq: Generates synthetic FASTQ records free of se-
 243 quences defined by regex patterns, or containing spiked sequences based
 244 on regex patterns.” <https://github.com/Rbfinch/spikeq>.

245 Free Software Foundation. 2023. *GNU Grep 3.11*. Free Software Foundation.
 246 <https://www.gnu.org/software/grep/manual/grep.html>.

247 ———. 2024. *GAWK: Effective AWK Programming: A User’s Guide for GNU*
 248 *Awk, for the 5.3.1*. Free Software Foundation. [https://www.gnu.org/softwa](https://www.gnu.org/software/gawk/manual/gawk.html)
 249 [re/gawk/manual/gawk.html](https://www.gnu.org/software/gawk/manual/gawk.html).

250 Gallant et al. 2025a. “rebar.” <https://github.com/BurntSushi/rebar>.

251 ——— et al. 2025b. “regex.” <https://github.com/rust-lang/regex>.

252 Gallant, Andrew. 2025. “Ripgrep: Recursively Search the Current Directory for
 253 Lines Matching a Pattern.” <https://github.com/BurntSushi/ripgrep>.

254 Hodgman, T. Charles. 2000. “A Historical Perspective on Gene/Protein Func-
 255 tional Assignment.” *Bioinformatics* 16 (1): 10–15.

256 Katz, Kenneth, Oleg Shutov, Richard Lapoint, Michael Kimelman, J Rodney
 257 Brister, and Christopher O’Sullivan. 2022. “The Sequence Read Archive:
 258 A Decade More of Explosive Growth.” *Nucleic Acids Research* 50 (D1):
 259 D387–90.

260 Kleene, SC. 1951. "Representation of Events in Nerve Nets and Finite Au-
 261 tomata." *CE Shannon and J. McCarthy*.

262 Martinez-Porchas, Marcel, Enrique Villalpando-Canchola, Luis Enrique Ortiz
 263 Suarez, and Francisco Vargas-Albores. 2017. "How Conserved Are the
 264 Conserved 16S-rRNA Regions?" *PeerJ* 5: e3036.

265 Merrett, James E, Monica Nolan, Leon Hartman, Nijoy John, Brianna Flynn,
 266 Louise Baker, Christelle Schang, et al. 2024. "Highly Sensitive Wastewater
 267 Surveillance of SARS-CoV-2 Variants by Targeted Next-Generation Ampli-
 268 con Sequencing Provides Early Warning of Incursion in Victoria, Australia."
 269 *Applied and Environmental Microbiology* 90 (8): e01497–23.

270 Mutiple. n.d. "Mimalloc: A Rust Wrapper over Microsoft's MiMalloc Memory
 271 Allocator."

272 Schlegel, Markus, and Adrian Seyboldt. 2025. "seq_io: FASTA and FASTQ
 273 parsing and writing in Rust." https://github.com/markschl/seq_io.

274 Sherry, Stephen, Chunlin Xiao, Kenneth Durbrow, Michael Kimelman, Kurt
 275 Rodarmer, Martin Shumway, and Eugene Yaschenko. 2012. "NCBI Sra
 276 Toolkit Technology for Next Generation Sequence Data." In *Plant and*
 277 *Animal Genome XX Conference (January 14-18, 2012)*. *Plant and Animal*
 278 *Genome*.

279 Sims, Natalie, and Barbara Kasprzyk-Hordern. 2020. "Future Perspectives
 280 of Wastewater-Based Epidemiology: Monitoring Infectious Disease Spread
 281 and Resistance to the Community Level." *Environment International* 139:
 282 105689.

283 Valdivia-Granda, Willy A. 2012. "Biodefense Oriented Genomic-Based
 284 Pathogen Classification Systems: Challenges and Opportunities." *Journal*
 285 *of Bioterrorism & Biodefense* 3 (1): 1000113.

286 Xylogiannopoulos, Konstantinos F. 2021. "Pattern Detection in Multiple

287 Genome Sequences with Applications: The Case of All SARS-CoV-2
288 Complete Variants.” *bioRxiv*, 2021–04.